

Visionary

12+3 · NO Discord / Drive content

I-JEPA = static-image birth certificate of Meta world-model line.

Actionable robotics = V-JEPA 2-AC. Hybrid: latent rollout + FM motor.

Spine: Predict embeddings, not pixels · multi-block + EMA + narrow predictor · JEA|MAE|I-JEPA ·
→V-JEPA→2-AC (cite 2-AC for robots) · 54.2 / 66.9»40.7 / ~10× / 72.4

Explicitly excluded: Discord bots, Drive sync plans, chat ops — research agenda only

BEHAVIOR reality — who won 2025

Winners $\neq$ JEPA stack

WINNERS RIDE

$\pi_{0.5}$ + flow-matching action experts + SigLIP / PaliGemma — a generative-action world-interface, not a JEPA latent dynamics model.

JEPA LINE = PARALLEL BET

Plan by optimizing actions so predicted latents match a goal-image embedding — almost no task demos (V-JEPA 2-AC recipe).

Do not pitch I-JEPA as the 2026 BEHAVIOR default trunk — pitch it as the WM parallel track.

Hybrid recipe — latent long + FM short

Best of both world-model bets

LONG-HORIZON

V-JEPA-style latent rollout

Feasibility / imagination / stage planning in embedding space — avoid wasting capacity on pixel video noise.

SHORT-HORIZON

FM / $\pi_{0.5}$ motor expert

Chunked flow-matching execution on SigLIP tokens — where BEHAVIOR winners already win contact-rich control.

Why hybrid for household long-horizon

Pure BC compounds error over stages; pure pixel video models waste capacity. Latent plan + FM execute is the actionable Visionary bet.

Consistency triangle → 2026 bets

Three trunks, three jobs

SigLIP / DINOv2

VLA default trunk
 $\pi_{0.5}$ / OpenVLA

Ship today

MAE-class

FT-heavy / dense
texture affordance

Ablation candidate

I-JEPA→2-AC

Latent plan /
image-goal CEM

WM parallel bet

Visionary rule: ablate JEPA features on BEHAVIOR/DROID stills (Empiricist E8) before any encoder FT politics.

Follow-up research agenda

Five tickets after this paper club

1. Action-conditioned JEPA

Post-train predictor on OXE / DROID / BEHAVIOR-1K video — compare sample efficiency to π_0 -style BC.

2. Goal-image CEM

Port V-JEPA 2-AC CEM planning into BEHAVIOR scenes; measure SR vs language-goal VLAs.

3. Uncertainty predictors

Ensemble / quantile heads for contact-rich failure anticipation.

4. Cross-embodied latent

Camera + proprio tokens; embodiment-agnostic predictor.

5. Affordance bake-off

Contrastive-vs-JEPA probes on BEHAVIOR affordances — not ImageNet-1K.

Action-conditioned path — OXE / DROID

Cite 2-AC as the existence proof

CLEAR existence proof

V-JEPA 2-AC: action-free JEPA on 1M+ hours video → action-conditioned post-train on <62 h DROID → zero-shot Franka pick-place with image goals / CEM — no env-specific reward.

Visionary next step

Repeat the recipe on BEHAVIOR-1K / OXE household slices. Metric: stages completed under OOD kitchens vs $\pi_{0.5}$ BC baseline at matched demo count.

Not in this 2023 PDF — do not attribute Franka SR to Assran et al. CVPR 2023.

Fig. 6 — predictor already imagines parts

Birth certificate of latent imagination




Figure 6. **Visualization of I-JEPA predictor representations.** For each image: first column contains the original image; second column contains the context image, which is processed by a pretrained I-JEPA ViT-H/14 encoder. Green bounding boxes in subsequent columns contain samples from a generative model decoding the output of the pretrained I-JEPA predictor, which is conditioned on positional mask tokens corresponding to the location of the green bounding box. Qualities that are common across samples represent information that contained is in the I-JEPA prediction. The I-JEPA predictor is correctly capturing positional uncertainty and producing high-level object parts with a correct pose (e.g., the back of the bird and top of a car). Qualities that vary across samples represent information that is not contained in the representation. In this case, the I-JEPA predictor discards the precise low-level details as well as background information.

Targets			Context		
Mask	Type	Err	Type	Avg. Ratio*	Top-1

RCDM decodes show pose/parts with uncertainty — the philosophical seed of goal-image latent planning.

New apps — beyond Discord / Drive

Research + robot products only

Zero-shot rearrange

Image goals on new kitchens via 2-AC recipe

Safety monitors

Predictor residual as surprise / anomaly

Sim-free WM pretrain

Self-supervised household world models

Education

Teach 'predictive not reconstructive' AMI slogan

Affordance probes

JEPA vs DINO on contact-rich stills

Stage imagination

Latent rollout for BEHAVIOR long-horizon stages

Anti-patterns

What Visionary refuses

✗ Pitching I-JEPA as drop-in OpenPI / $\pi_{0.5}$ vision tower without E8 probe card.

✗ Claiming BEHAVIOR q-score gains from citing 79.3% ImageNet linear.

✗ Putting Franka SR next to Assran 2023 without 2-AC citation.

✗ Shipping Discord bots / Drive sync as the 'follow-up plan'.

✗ Retraining paper-scale ViT-H/14 as a student milestone (E10 = politics).

✗ Treating JEPA and FM experts as mutually exclusive — hybrid is the bet.

One-liner — tattoo this sentence

KanBot / exam ready

I-JEPA is the static-image birth certificate of Meta's world-model line; the actionable robotics paper is V-JEPA 2 / 2-AC — but you cannot understand that line without this one.

Companion: predict embeddings, not pixels · multi-block 54.2 · emb 66.9>>40.7 · cite 2-AC for robots.

Exam bite — Visionary

Parallel bet, not default trunk

BEHAVIOR winners use $\pi_{0.5}$ +FM+SigLIP. JEPA is the parallel latent-plan bet (V-JEPA 2-AC). Hybrid = latent rollout for long-horizon + FM for short-horizon motor. Follow-ups: action-conditioned JEPA on OXE/DROID, goal-image CEM in BEHAVIOR, uncertainty predictors. No Discord/Drive. Birth certificate = I-JEPA; actionable robots = 2-AC.

Numbers tattoo — pin these

Shared across all five seats

79.3%

ViT-H/14 linear

73.3%

H/14 · 1% labels

81.1% / 77.3%

H/16@448 lin / 1%

54.2

multi-block (T6)

66.9 >> 40.7

emb >> pixels (T7)

72.4

Clevr Dist (T4)

~10×

cheaper than MAE H/14

<72 h

16×A100 H/14

<1200

GPU-hours H/14

87.1 vs 87.8

FT H/16@448 vs MAE

no view augs

single-view JEPA

2-AC not here

robots → cite 2506.09985

Visionary — close

Takeaways + exam bite

1. Winners: $\pi_{0.5}$ +FM+SigLIP — JEPA line is the parallel latent-plan bet.
2. Hybrid: V-JEPA-style latent rollout (long) + FM motor (short).
3. Follow-ups: action-JEPA on OXE/DROID; goal-image CEM; uncertainty heads.
4. One-liner: birth certificate = I-JEPA; actionable robotics = V-JEPA 2-AC.

EXAM BITE

Parallel WM bet + hybrid recipe; cite 2-AC for robots; no Discord/Drive.